

# Neural Networks II

Week 15

# Applications

- Natural language processing (NLP)
- Speech recognition and synthesis
- Image recognition
- Self-driving cars
- Customer experience, healthcare, and robotics

# Linear Regression

- Key foundation for deep learning
  - A linear model
  - Relationship between two or more variables
  - Predict dependent variable based on independent variable
  - Slope and intercept models the relationship

y	Dependent Variable
x	Independent Variable
a	Slope
b	Intercept

$$y = ax + b$$

$$y = a_1x_1 + a_2x_2 \dots + a_nx_n + b$$

# Building a linear regression model

- Find values for slope and intercept
- Use known values of x and y (multiple samples)
- Multiple independent variables make it complex

$$5 = 2a + b, 9 = 4a + b$$

$$b = 5 - 2a$$

$$9 = 4a + 5 - 2a$$

$$a = (9-5)/2 = 2$$

$$b = 5 - 2*2 = 1$$

# Logistic regression

- A binary model
- Relationship between two or more variables
- Output is 0 or 1

<b>y</b>	Dependent Variable
<b>x</b>	Independent Variable
<b>a</b>	Slope
<b>b</b>	Intercept
<b>f</b>	Activation Function

$$y = f(ax + b)$$

$$y = f(a_1x_1 + a_2x_2 \dots + a_nx_n + b)$$

# An analogy for Deep Learning

- Complex and iterative process
- Starts with random initialization and works towards the right values by trial and error


$$10 = 3a + b$$

Trial	a	b	$3a + b$	Error
1	1	1	4	6
2	4	3	15	-5
3	2	2	8	2
4	3	2	11	-1
5	2	3	9	1
6	2	4	10	0



# The Perceptron

- An algorithm for supervised learning for binary classification
- Resembles a cell in the human brain
- A single cell neural network
- Based on logistic regression



# Perceptron formula

$$y = f(a_1x_1 + a_2x_2 \dots + a_nx_n + b)$$



$$y = f(w_1x_1 + w_2x_2 \dots + w_nx_n + b)$$

<b>w</b>	Weight
<b>b</b>	Bias
<b>f</b>	Activation Function


## **Activation Function Example**

*1 if value > 0*

*0 otherwise*




# Perceptron



# Artificial Neural Network

- A network of perceptron
- Perceptrons are called nodes in the neural network
- Nodes organized as layers
- Each node has weights, biases and activation functions
- Each node is connected to all nodes in the next layer

# Artificial Neural Network



# How ANN works for prediction?

- Inputs (independent variables) are sent from the input layer
- Inputs passed on to the nodes in the next hidden layer
- Each node computes its output based on its weights, biases, and activation functions
- Node output is then passed on as inputs to the next layer


# Training an ANN

**A model is represented by parameters and hyperparameters.**

- Weights, biases, nodes, layers, layers, etc.

**Training a model means determining the best possible values for the parameters and hyperparameters that maximize accuracy.**

**Inputs, weights, and biases might be n-dimensional arrays.**



# Recall the Linear Regression Analogy

Weight		Bias			
Trial	a	b	$3a + b$	Error	
1	1	1	4	6	
2	4	3	15	-5	
3	2	2	8	2	
4	3	2	11	-1	
5	2	3	9	1	
6	2	4	10	0	



# Training Process

- Use training data (known values of inputs and outputs)
- Create network architecture with intuition
- Start with random values for weights and biases
- Minimize error in predicting known outputs from inputs
- Adjust weights and biases until error is minimized
- Improve model by adjusting layers, node counts and other hyper parameters



# The Input Layer

**The input to deep learning is a vector of numeric values**

**A vector is a tuple of one or more values**

- E.g., ( 1.0, 2.1, 3.5)

**Defined usually as a NumPy array**

**Represents the feature variables for prediction**



# Samples and Features

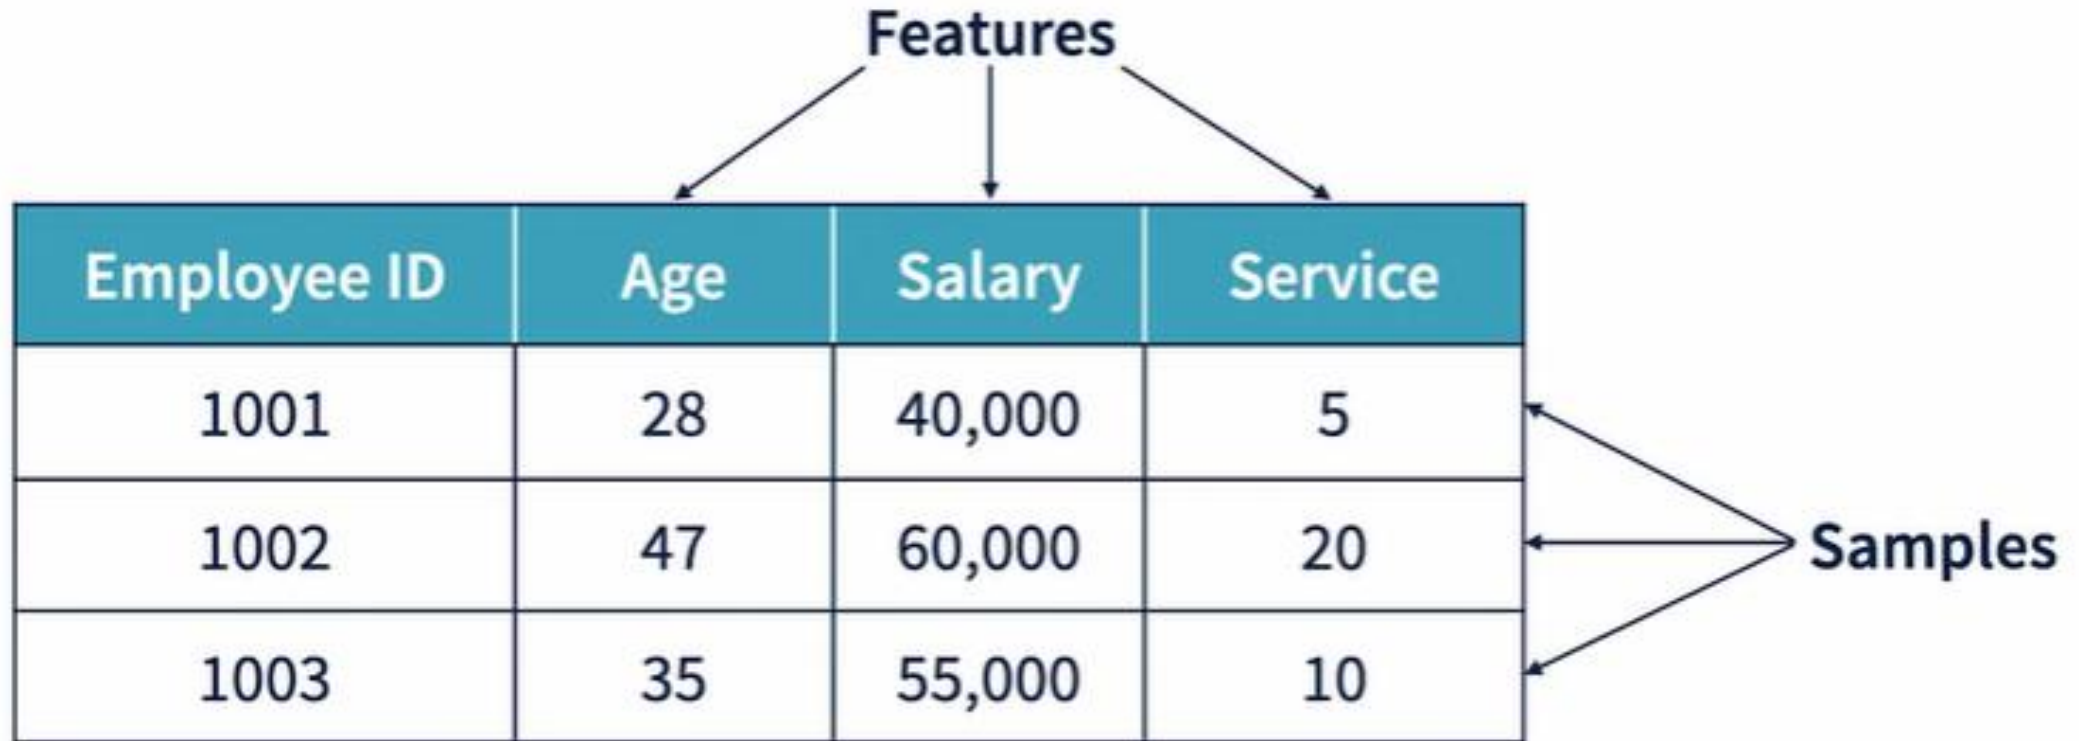


Diagram illustrating the relationship between Features and Samples in a dataset.

The word **Features** is positioned above the table, with arrows pointing to the columns: **Age**, **Salary**, and **Service**.

The word **Samples** is positioned to the right of the table, with arrows pointing to the rows (1001, 1002, and 1003).

Employee ID	Age	Salary	Service
1001	28	40,000	5
1002	47	60,000	20
1003	35	55,000	10

# Input Pre-Processing

- Features need to be converted to numeric representations

Input Type	Preprocessing Needed
Numeric	Centering and scaling
Categorical	Integer encoding, one-hot encoding
Text	TF-IDF, embeddings
Image	Pixels – RGB representation
Speech	Time series of numbers

# Example

**Raw Data**

Age	Salary	Service
28	40,000	5
47	60,000	20
35	55,000	10

**Centered and Scaled**


x1	x2	x3
-1.10	-1.37	-1.07
1.32	0.98	1.34
-0.21	0.39	-0.27

**Transposed**

x1	-1.1	1.32	-0.21
x2	-1.37	0.98	0.39
x3	-1.07	1.34	-0.27


# Hidden Layers

- An ANN can have one or more hidden layers
- Each hidden layer can have one or more nodes (or perceptron) (count in  $2^n$ )
- A neural network architecture is defined by the number of layers and nodes



# Inputs and Outputs

- The output of each node in the previous layer becomes the input for each node in the current layer (fully connected)
- Each node produces one output that is forwarded to the next layer






# How to determine hidden layer architecture?

Each node *learns* something about the feature-target relationship

More nodes and layers mean


- Web applications on AWS
- Analytics on GCP

Architecture decided by experimentation




# Weights and biases

- Weights and biases represent the trainable parameters in an ANN
- Numeric values
- Each input for each node has a weight associated with it
- Each node has a bias associated with it



# Weights and biases for a layer




Layer	Inputs	Nodes	Weights	Biases
HL 1	3	4	12	4
HL 2	4	5	20	5
HL 3	5	3	15	3
Output	3	2	6	2
<b>Total</b>			<b>53</b>	<b>14</b>



# Computing Outputs

- This example shows the computation for hidden layer two which has four inputs and five nodes.




$$\begin{bmatrix} w_{11} & w_{21} & w_{31} & w_{41} \\ w_{12} & w_{22} & w_{32} & w_{42} \\ w_{13} & w_{23} & w_{34} & w_{43} \\ w_{14} & w_{24} & w_{34} & w_{44} \\ w_{15} & w_{25} & w_{35} & w_{45} \end{bmatrix} \star \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \\ b_5 \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \\ y_5 \end{bmatrix}$$

$$W^T X + B = Y$$

# Activation functions

- Determines which nodes propagate information to the next layer
- Filters and normalizes
- Converts output to nonlinear
- Critical in learning patterns




# Popular Activation Functions

Activation Function	Output
Sigmoid	0 to 1
Tanh	-1 to +1
Rectified Linear Unit (ReLU)	0 if $x < 0$ ; $x$ otherwise
Softmax	Vector of probabilities, with sum=1

Choice depends on problem and experimentation.

# The Output Layer


- One layer of output, produces desired Y
- Has its own weights and biases
- Softmax activation used for classification problems
- May need postprocessing to convert to business values



# Output Layer size

## Size depends on the problem

- 1 for binary classification
- $n$  for a  $n$ -class classification
- 1 for regression problems
- Vary based on other problem domains





# Training the network – data pre-processing

	Sample 1	Sample 2	Sample 3	Sample 4
Feature 1	x11	x21	x31	x41
Feature 2	x12	x22	x32	x42
Feature 3	x13	x23	x34	x43
Feature 4	x14	x24	x34	x44
Feature 5	x15	x25	x35	x45
Target	y1	y2	y3	y4

# Split input

- Training set: Used to fit the parameters
- Validation set: Used for model selection/tuning
- Test set: Used to measure the final model performance
- Usual split: 80:10:10

# Select values for the model

## Select values for the model

- Layers and nodes in the layer, activation functions
- Hyper parameters

## Selection criteria

- Initial selection based on intuition/reference
- Adjustment based on results






# Initialize weights

- All weights and bias parameters need to be initialized to some value before we start training
- Zero initialization: Initialize to zeros, not recommended
- Random initialization: Initialize to random values from a standard normal distribution (mean = 0, SD = 1)

# Forward Propagation

	Sample 1	Sample 2	Sample 3	Sample 4
Feature 1	x11	x21	x31	x41
Feature 2	x12	x22	x32	x42
Feature 3	x13	x23	x34	x43
Target	y1	y2	y3	y4
Prediction	$\hat{y}1$	$\hat{y}2$	$\hat{y}3$	$\hat{y}4$


# Forward propagation: 1 Sample



# Forward propagation: All Samples

- Send each sample through the neural network and obtain the value of  $\hat{y}$
- Repeat for all samples and collect a set of  $\hat{y}$
- Compare the values of  $\hat{y}$  to  $y$  to obtain error rates

# Measuring Accuracy and Error





# Loss and Cost function

- A loss function measures the prediction error for a single sample
- A cost function measures the error across a set of samples
- Popular Cost Functions

Cost Functions	Applications
Mean Square Error ( MSE )	Regression
Root Mean Square Error ( RMSE )	Regression
Binary Cross Entropy	Binary classification
Categorical Cross Entropy	Multi-class classification

# Measuring Accuracy

- Send a set of samples through the ANN and predict outcome
- Estimate the prediction error between the predicted outcome and expected outcome using a cost function
- Use back propagation to adjust weights based on the error value

# Back Propagation

- Each node in a neural network contributes to the overall error in prediction (differing contributions)
- A node's contribution is driven by its weights and bias
- Weights and biases need to be adjusted to lower the error contribution by each node




# How does it work?

- It works in reverse of the forward propagation.
- Start from the output layer
- Compute a delta value based on the error found
- Apply the delta to adjust the weights and biases in the layer
- Derive a new error value
- Back propagate the new error to the previous layer and repeat

Input

Weights and Biases



# Gradient Descent

## Repeat the learning process.

- Forward propagation
- Estimate error
- Backward propagate
- Adjust weights and biases



# Batch

- A set of samples sent through ANN in a single pass
- The training data set can be divided into one or more batches
- Training data is sent to the ANN one batch at a time
- Cost estimated and parameters updated one batch at a time
- Batch gradient descent
  - Batch size = training set size
- Mini-batch gradient descent
  - Batch size  $<$  training set size
- Typical batch sizes are 32, 64, 128, etc.

# Epoch

- The number of times the entire training set is sent through the ANN
- An epoch has one or more batches
- The training process completes when all epoch is complete
- Epoch sizes can be higher to achieve better accuracy

# Epoch and Batch Example

- Training set size = 1000, batch size = 128, epoch = 50
- Batches per epoch =  $\text{ceil} ( 1000 / 128 ) = 8$
- Total iterations (passes) through ANN =  $8 * 50 = 400$
- Batch size and epoch are hyperparameters that can be tuned to improve model accuracy



# Validation and Testing

- Validation

- During learning, the predictions are obtained for the same data that is used to train the parameters (weights and biases)
- After each epoch and corresponding parameter updates, the model can be used to predict for the validation data set
- Accuracy and/or loss can be measured and investigated
- Model can be fine-tuned and learning process repeated based on results.

- Evaluation

- After all fine-tuning is completed and final model obtained, the test data set can be used to evaluate the model
- Results obtained with test data is used to measure the performance of the model

# Summary: An ANN Model

## Parameters

- Weights
- Biases

## Hyperparameters

- Number of layers, nodes in each layer, activation function
- Cost functions, learning rate, optimizers
- Batch size, epoch



# Summary: Prediction Process

**Preprocess and prepare inputs**

**Pass inputs to the first layer**

- Compute  $Y$  using weights, biases, activation
- Pass to the next layer

**Repeat process until output layer**

**Postprocess output for predictions**

